

Connections AI Solver

Studienarbeit

for the Informatik Department at the
DHBW Heidenheim

by
Ziyi Hong

TODO. TODO 2026

Processing period	January – March 2026, July – September 2026
Student ID	2694933
Course	Allgemeine Informatik
Evaluator	Tobias Heß

Declaration of Independence

I hereby declare that I have written my study on the topic “*Connections AI Solver*” independently and that I have not used any sources or aids other than the indicated ones. I also declare that the submitted electronic version matches the printed version.

MyPlace, July 30, 2026

Place, Date

Signature

Abstract

Contents

List of Tables	VI
1 Introduction	1
1.1 Structure	1
2 Foundations	2
2.1 The NYT Connections Puzzle	2
2.2 Relation Types and the Structure of Difficulty	3
2.3 Challenges for Automated Solvers	4
2.4 Small and Resource-Constrained Language Models as Solving Agents . . .	5
2.5 Text Embeddings and Clustering for Lexical Grouping	6
2.6 Prompting Strategies for LLM Reasoning	6
2.7 Related Work on NYT Connections Solvers	7
3 Related Work	10
3.1 Prior Work on Connections Solving	10
4 Experimental Design and Implementation	12
5 Evaluation and Discussion	13
5.1 Experiments	13
6 Conclusion and Future Work	14
Bibliography	IV

CoT Chain-of-Thought

List of Tables

2.1	Relation types in NYT Connections, with illustrative examples.	3
-----	--	---

1 Introduction

NYT Connections asks players to sort sixteen words into four groups of four. Each group carries one of four difficulty tiers, ranging from yellow (easiest) to purple (hardest). Published daily, the puzzle has become a popular benchmark for testing the reasoning abilities of language models. Two challenges make Connections difficult to solve.

The first challenge is lexical. Solving Connections requires several distinct kinds of language understanding at once. Some groups rest on straightforward synonymy. Others depend on shared themes, wordplay, or knowledge of pop culture and named entities. A single puzzle typically mixes several of these relation types, and is deliberately constructed so that some words could plausibly belong to more than one candidate group under a superficial reading, an effect commonly referred to as a red herring.

The second challenge is structural. Because a word cannot be reused across groups, correctly identifying even one group depends not only on recognizing a plausible shared category among four words in isolation, but also on the arrangement of the remaining words. A locally plausible grouping can still be wrong if it leaves the remaining words unable to form additional valid groups of four.

In previous studies, the first challenge has been widely discussed, through embedding-based clustering or large language models to propose or judge candidate groupings Section 3.1. This thesis evaluates how far such approaches can be pushed, including several forms of knowledge injection beyond raw embeddings, and finds that they consistently fall short in similar ways (note: to be written, do the reference later!!!). The second challenge, ensuring that all sixteen words are partitioned simultaneously and exclusively, has received comparatively little attention on its own. Taking inspiration from solving Wordle algorithmically with information theory [1], this thesis treats the puzzle’s interactive structure, submitting a guess, observing feedback, and revising, as central to addressing this second challenge, and asks how much of the puzzle’s difficulty it can resolve by algorithm.

1.1 Structure

2 Foundations

This chapter establishes the conceptual and technical foundation for this thesis. Section 2.1 introduces the NYT Connections puzzle and its rules. Section 2.2 develops a taxonomy of the semantic relations underlying puzzle categories and links it to the puzzle’s difficulty structure. Section 2.3 discusses why the puzzle is difficult for automated solvers in general. Section 2.4 motivates the focus on small, locally deployable language models. Sections 2.5 and 2.6 cover the two technical building blocks of the proposed solver – text embeddings/clustering and LLM prompting strategies. Section 2.7 reviews prior solver designs for Connections and positions this thesis relative to them.

2.1 The NYT Connections Puzzle

NYT Connections is a daily word puzzle in which players must partition a set of sixteen words into four disjoint groups of four. Each correct group is assigned one of four color-coded difficulty labels – yellow (easiest), green, blue, and purple (hardest) – reflecting the puzzle’s intended difficulty in recognizing the underlying grouping relation.

Two evaluation protocols are relevant for this thesis and are used interchangeably across the experiments described in ??¹:

- An **interactive setting** with a limited number of attempts (four in the original game) and “one-away” feedback: after each submitted group, the player learns whether the group was fully correct, wrong, or contained exactly three correct words (“one away”).
- An **all-at-once setting** that requires predicting the complete 4×4 partition in a single submission, without intermediate feedback.

Solver communities have also adopted self-imposed variants to increase difficulty, such as attempting the purple category first (“purple-first”), which reverses the typical order of play by prioritizing the most abstract relation before the easier categories are identified. This practice illustrates that even human solvers treat the difficulty ordering yel-

¹Adjust the chapter label `ch:methodology` once your methodology chapter exists.

low → green → blue → purple as a meaningful, exploitable structure rather than an arbitrary label – a point taken up again in Section 2.2.

2.2 Relation Types and the Structure of Difficulty

The categories in a Connections puzzle are not homogeneous in the kind of linguistic or world knowledge they require. Following the taxonomy summarized in Table 2.1, four broad relation types can be distinguished, with a fifth “mixed” category covering cases that combine more than one type.

Table 2.1: Relation types in NYT Connections, with illustrative examples.

Relation type	Example category and words
Semantic similarity	<i>FURTHEST POINT</i> : END, EXTREME, OPPOSITE, POLE [nyt-puzzle-2026-01-05]
Semantic relatedness	<i>THINGS YOU CAN DO ON SOCIAL MEDIA</i> : COMMENT, LIKE, LURK, POST [nyt-puzzle-2026-01-05]
Form- or rule-based relations	<i>ART MOVEMENTS, WITH “-ISM”</i> : BRUTAL, IMPRESSION, MANNER, REAL [nyt-puzzle-2026-01-05]
Referential relations	<i>RIHANNA #1 HITS</i> : DIAMONDS, SOS, UMBRELLA, WORK [nyt-puzzle-2026-01-03]
Mixed categories	<i>CAR BRAND HOMOPHONES</i> : INFINITY, MINNIE, OPAL, OUTIE [nyt-puzzle-2025-12-29]

Semantic similarity relations rest on near-synonymy or shared hypernymy (e.g., “types of flowers”); semantic relatedness relations hold between words that are not strictly similar but co-occur within a common theme or script; form- or rule-based relations are determined by surface regularities such as spelling, morphology, or conventional symbol meanings; referential relations require external world knowledge about named entities or factual lists. The hardest (purple) categories draw disproportionately on the less standard relation types – word and letter patterns such as rhymes, homophones, anagrams, or palindromes, unmarked fill-in-the-blank phrases, and pop-culture or proper-noun knowledge –

and frequently combine more than one relation type at once [nyt-tips-tricks-2023].

Working hypothesis adopted in this thesis. The color-coded difficulty ordering correlates, at least approximately, with this relation-type taxonomy: yellow and green categories tend toward semantic similarity and relatedness, which are directly addressable by distributional/embedding-based methods, whereas blue and purple categories skew toward form-based, referential, and mixed relations, which embeddings alone are ill-suited to capture and which instead require either symbolic pattern matching or LLM-internal world knowledge. This mapping is not drawn from a specific external source but is proposed here as an explanatory lens; it is revisited in ??² as the basis for the per-tier error analysis.

2.3 Challenges for Automated Solvers

The relation-type diversity described above already indicates why Connections is difficult for automated approaches: a single puzzle may combine semantic similarity, semantic relatedness, form- or rule-based cues, and referential knowledge. Consequently, the task spans multiple NLP subproblems simultaneously – distributional semantics and sense disambiguation for similarity/relatedness categories, pattern matching for form-based categories, and retrieval of external, often culturally specific knowledge for referential categories – so that no single notion of similarity is sufficient across all groups.

The puzzle is also difficult by design. Categories frequently overlap in plausibility, creating red herrings: words that appear to fit one candidate grouping but ultimately belong elsewhere. Context clues are consequently critical for disambiguating individual words, and because the four groups must jointly partition all sixteen words, early (incorrect) commitments constrain what remains achievable for the rest of the puzzle [nyt-tips-tricks-2023].

Loredo Lopez, McDonald, and Emami [loredolopez2025] explicitly design their Connections benchmark to penalize fast, superficial-cue-driven (“System-1”) guesses. This framing aligns with the shortcut-learning perspective: models can succeed by exploiting easy-to-use regularities rather than recovering the intended global grouping structure, a

²Adjust once your results chapter exists.

behavior expected whenever the data distribution offers shortcut opportunities and the evaluation objective rewards whatever features are most locally predictive [geirhos2020]. Connections’ daily release cadence additionally supports robust, leakage-resistant model comparison, since each day’s puzzle is unseen at evaluation time relative to a model’s training cutoff [loredolopez2025].

2.4 Small and Resource-Constrained Language Models as Solving Agents

Most reported LLM results on Connections-style benchmarks use large, cloud-hosted models [toddetal2024, loredolopez2025, zhangetal-connections]. This thesis instead centers on small, locally executable models in the 7–8B parameter range (e.g., Mistral 7B, Llama 3.1 8B), run through a local inference server (Ollama) rather than a hosted API. Three considerations motivate this focus:

1. **Deployment constraints.** Local execution removes dependence on external API availability, per-token cost, and data egress, which is relevant whenever a solver is meant to run on constrained or offline infrastructure.
2. **The capability gap as the actual object of study.** Because small models trail larger ones on complex reasoning tasks, they expose failure modes – shortcut reliance, weak constraint tracking across a 16-word global partition, limited world knowledge for referential categories – more readily than large models do. This makes small models a more informative testbed for evaluating whether *external scaffolding* (embeddings, clustering, structured verification) can substitute for missing model-internal capability, which is the central question this thesis investigates.
3. **Comparability to prior work.** Loredolopez, McDonald, and Emami [loredolopez2025] already evaluate multiple LLMs of varying size under different prompting schemes on Connections, which provides a point of comparison for the relative performance of the small models used here (Section 2.7).

2.5 Text Embeddings and Clustering for Lexical Grouping

At the representation level, Connections is closely tied to text embeddings, semantic similarity, and clustering. Early work on distributed word representations introduced dense word vectors capturing lexical regularities through co-occurrence statistics [mikolov2013]. Subsequent surveys trace the shift from such static word vectors to contextualized representations, including sentence-level embeddings optimized for semantic similarity [chandrasekaran2021]. transformer-based embedding models are now the standard approach for encoding context-dependent lexical and semantic relations [incitti2023], and a broader survey perspective summarizes representation choices and their implications for similarity modeling across NLP tasks more generally [patil2023].

Classical clustering algorithms – k-means, k-means++, and DBSCAN, together with keyword-clustering techniques – group items based on distances within such embedding spaces [ester1996dbscan, arthur2007kmeanspp]. More recent work shows that LLM-derived embeddings can improve resulting cluster structure, and that clustering itself can be reframed as an in-context classification problem solved directly by an LLM rather than by a separate distance-based algorithm [huang2025textclustering, petukhova2025].

This thesis uses SentenceTransformer-style encoders (specifically BAAI/bge-large-en-v1.5, alongside comparison runs with other encoder families such as the E5 series and MiniLM) as the embedding backbone for the non-LLM baseline and for the clustering pre-processing stage of the hybrid pipeline (??). The choice of embedding family is treated as an experimental variable rather than fixed a priori, following the observation in Section 2.7 that prior work on Connections instantiates broadly comparable solver structures with substantially different embedding families, without directly comparing them under a controlled setup.

2.6 Prompting Strategies for LLM Reasoning

Chain-of-thought (CoT) prompting elicits intermediate reasoning steps from an LLM before it commits to a final answer and has been shown to improve performance on complex, multi-step tasks [wei2022cot]. Prompting-technique reference guides collect CoT and related variants in a practitioner-oriented form [promptguide-cot].

The solver pipelines evaluated in this thesis instantiate several points along this spectrum: a zero-shot condition, a few-shot condition using a single worked example, and a CoT condition that decomposes reasoning into an explicit analyze-group-verify sequence (see `system_prompt()` in the implementation). A self-consistency variant, which samples multiple reasoning traces and aggregates over them rather than committing to a single greedy completion, is evaluated as an additional condition. Loredo Lopez, McDonald, and Emami [loredolopez2025] similarly contrast direct input-output prompting against CoT and self-consistency variants on Connections specifically, which this thesis uses as a point of reference in Section 2.7 and ??.

2.7 Related Work on NYT Connections Solvers

Three prior studies directly inform the solver design used in this thesis.

Todd et al. [toddetal2024] study Connections using both embedding-based heuristics and prompted LLMs. Their non-LLM baseline encodes words with SentenceTransformer-style models, exhaustively scores all $\binom{16}{4}$ candidate four-word groups by within-group cosine-similarity cohesion, and constructs a solution iteratively; they also consider an all-at-once variant that ranks complete 4×4 partitions by summed cohesion.

Loredo Lopez, McDonald, and Emami [loredolopez2025] propose a “System-1-like” heuristic baseline using Multilingual-E5-Large-Instruct embeddings, scoring candidate groups with a discriminative function

$$S = G - P, \tag{2.1}$$

where G aggregates multiple within-group cohesion signals and P penalizes groups that are insufficiently distinctive from the remaining words. Full solutions are assembled via beam search. They evaluate multiple LLMs under direct, chain-of-thought, and self-consistency prompting schemes.

Correspondence to this project’s implementation. The `score_group()` function in `herustic_solver.py` computes

$$G = 0.4 \cdot I + 0.3 \cdot s + 0.3 \cdot V$$

(a weighted combination of negative inertia, minimum pairwise cosine similarity, and a mean/variance cohesion term) and a final score $S = G - P$, assembled via a beam-search procedure over group choices – structurally identical to the scoring function in Equation 2.1 as described by Loredolo Lopez et al. [loredolopez2025]. **This correspondence must be stated explicitly in the methodology chapter** as the basis for the reproduced baseline, both to satisfy the “solides Fundament” principle of the DHBW guideline and to avoid any appearance of an unattributed structural borrowing.

Zhang, Jiang, and Ye [zhangetal-connections] frame the task as equal-size constrained clustering and apply Balanced k-means on static word embeddings (GloVe/-word2vec), with an alternative KMeans-plus-rebalancing baseline to enforce the four-groups-of-four partition. They additionally explore LLM-based pipelines including in-context learning, parameter-efficient fine-tuning (QLoRA), rationale distillation from a larger teacher model, and feedback-driven refinement.

Synthesis and positioning of this thesis. These three studies converge on a common structural insight – that a Connections solver decomposes into (a) a representation/scoring stage over candidate groups and (b) an assignment stage that resolves the equal-size global partition – while differing substantially in how each stage is instantiated: embedding family (E5 [loredolopez2025] vs. SentenceTransformer/MPNet [toddetal2024] vs. static word2vec/GloVe [zhangetal-connections]), scoring function, and assignment method (iterative greedy construction [toddetal2024], beam search [loredolopez2025], balanced/rebalanced k-means [zhangetal-connections]). None of the three directly compares embedding families under an otherwise fixed solver structure, and only Zhang, Jiang, and Ye [zhangetal-connections] treat the assignment stage as an explicitly constrained combinatorial problem – but resolve it with a clustering heuristic rather than an exact method.

This thesis addresses both gaps directly. First, it reproduces the [loredolopez2025] scoring function as a fixed non-LLM baseline and instantiates it with multiple embedding families under an otherwise controlled setup (??), extending the embedding-family comparison that [toddetal2024, loredolopez2025, zhangetal-connections] motivate but

do not themselves carry out. Second, building on the constrained-clustering framing of **[zhangetal-connections]**, it replaces the balanced/rebalanced clustering heuristic with exact and near-exact constrained-assignment methods (dynamic programming and integer linear programming via OR-Tools) applied to LLM- and embedding-scored candidate groups, combined with an explicit LLM-based trap/red-herring verification stage – directly addressing the red-herring challenge identified in Section 2.3 **[nyt-tips-tricks-2023]**. Third, unlike **[toddetal2024, loredolopez2025, zhangetal-connections]**, which evaluate primarily at the level of overall puzzle success, this thesis evaluates performance separately per difficulty tier (Section 2.2), which is necessary to assess the specific research question of whether the proposed hybrid pipeline improves solving accuracy on the easier (yellow/green) categories relative to the small-model baselines established in Section 2.4.

3 Related Work

This chapter provides an overview of prior work relevant to this thesis, including work that directly solves Connections using LLMs and work that inspired the methods used in this study. It first introduces prior work on solving Connections with AI. It then presents related work on the methods used, which serves as the knowledge base for this study.

3.1 Prior Work on Connections Solving

Current work on this topic mainly uses LLMs as the main solution for solving Connections. These studies compare how well different models solve Connections. Some use Connections as a benchmark to test LLM capability, while others explore different prompting techniques and methods. Todd et al. use three methods, namely a sentence-embedding baseline, a comparison between two LLMs (GPT-3.5-Turbo-1106 and GPT-4-1106-Preview), and Chain-of-Thought (CoT) prompting. Their experiments allow five incorrect and five invalid guesses per puzzle, which departs from the original game’s rules. GPT-4 Turbo with CoT achieves the highest success rate at 58.11%. The best-performing sentence-embedding baseline (MPNet) reaches 11.6%. LLMs are often stumped by categories that involve non-semantic word properties, abstract features, or context-dependent usage. They also report a finding relevant to this thesis. Puzzle order matters, since presenting puzzles in their original, grouped order achieves substantially higher accuracy than presenting them in randomized order [2].

A similar work from Samadarshi et al. benchmarked several LLMs against both novice and expert human players and creates a taxonomy of knowledge types required to solve connections. They used in total 438 Connections puzzles and their best performing models Claude Sonnet 3.5 achieved 18% success rate of fully solving the puzzles. They categorizes NYT Connections categories along two main dimensions, word form (covering morphology, orthography, phonology, and multiword expressions) and word meaning (covering semantic relations such as synonymy, polysemy, and hypernymy, along with associative and encyclopedic relations), with a further category for groups that combine both form-based and meaning-based cues. Results show that models performed comparably well on semantic relation categories but significantly worse on the other types.

3 Related Work

More recent work from Loredó Lopez, McDonald, and Emami compare three game configurations: single attempt, four tries without feedback (one away prompt) and four tries with feedback. Human solvers outperform LLM under all three configurations. The best performing model Claude 3.5 under full hints reaches 40.25%, while the human full-hints score reaches 60.67%. They also experiment on different prompting techniques. CoT and self-consistency prompting yield diminishing or even negative returns as puzzle difficulty increases. Simple input-output prompting sometimes outperforms these more elaborate reasoning strategies. Their heuristic (non-LLM) baseline, using Multilingual-E5-Large-Instruct embedding and beam search, reaches 13.25% accuracy.

Zhang, Jiang, and Ye compare in-context prompting of a large model (DBRX) against fine-tuning and distilling a smaller model (Mistral 7B) on solved puzzles. They found that fine-tuning a smaller model (Mistral 7B) directly on solved puzzles fails to capture the task itself, which underperformed even their clustering baseline. Their strongest result comes from Verbal Reinforcement Learning, in which a larger supervisor model critiques the fine-tuned model’s output in natural language across three rounds, roughly doubling the average per-puzzle match rate relative to single-pass prompting [5].

4 Experimental Design and Implementation

5 Evaluation and Discussion

5.1 Experiments

How do we build it: data, pipeline, modules.

- Pure embedding + clustering solver
- Pure LLM solver
- Embedding + LLM solver
- Improvement

6 Conclusion and Future Work

Bibliography

- [1] *Solving Wordle Using Information Theory*. Feb. 6, 2022. URL: <https://www.youtube.com/watch?v=v68zYyaEmEA> (visited on 07/30/2026).
- [2] Graham Todd et al. “Missed Connections: Lateral Thinking Puzzles for Large Language Models”. In: *2024 IEEE Conference on Games (CoG)*. 2024 IEEE Conference on Games (CoG). Aug. 2024, pp. 1–8. DOI: 10.1109/CoG60054.2024.10645557. URL: <https://ieeexplore.ieee.org/document/10645557> (visited on 01/05/2026).
- [3] Prisha Samadarshi et al. *Connecting the Dots: Evaluating Abstract Reasoning Capabilities of LLMs Using the New York Times Connections Word Game*. Oct. 14, 2024. DOI: 10.48550/arXiv.2406.11012. arXiv: 2406.11012 [cs.CL]. URL: <http://arxiv.org/abs/2406.11012> (visited on 07/02/2026). Pre-published.
- [4] Angel Yahir Loredó Lopez, Tyler McDonald, and Ali Emami. “NYT-Connections: A Deceptively Simple Text Classification Task That Stumps System-1 Thinkers”. In: *Proceedings of the 31st International Conference on Computational Linguistics*. COLING 2025. Ed. by Owen Rambow et al. Abu Dhabi, UAE: Association for Computational Linguistics, Jan. 2025, pp. 1952–1963. URL: <https://aclanthology.org/2025.coling-main.134/> (visited on 12/17/2025).
- [5] Emily Zhang, Yanan Jiang, and Peixuan Ye. *Learning Semantic Complexities of NYT Connections*. URL: <https://web.stanford.edu/class/cs224n/final-reports/256847963.pdf>.